

TUTORIAL

How To Set Up an Elasticsearch, Fluentd and Kibana (EFK) Logging Stack on Kubernetes

Logging Elasticsearch Kubernetes Solutions

By [Hanif Jetha](#)

Last Validated on March 30, 2020 • Originally Published on November 26, 2018 © 239.1k

🌐 English ▼

Introduction

When running multiple services and applications on a Kubernetes cluster, a centralized, cluster-level logging stack can help you quickly sort through and analyze the heavy volume of log data produced by your Pods. One popular centralized logging solution is the **Elasticsearch**, **Fluentd**, and **Kibana** (EFK) stack.

Elasticsearch is a real-time, distributed search engine which allows for full-text and structured search, as well as analytics. It is commonly used to index and search

SCROLL TO TOP

through large volumes of log data, but can also be used to search many different kinds of documents.

Elasticsearch is commonly deployed alongside **Kibana**, a powerful data visualization frontend and dashboard for Elasticsearch. Kibana allows you to explore your Elasticsearch log data through a web interface, and build dashboards and queries to quickly answer questions and gain insight into your Kubernetes applications.

In this tutorial we'll use **Fluentd** to collect, transform, and ship log data to the Elasticsearch backend. Fluentd is a popular open-source data collector that we'll set up on our Kubernetes nodes to tail container log files, filter and transform the log data, and deliver it to the Elasticsearch cluster, where it will be indexed and stored.

We'll begin by configuring and launching a scalable Elasticsearch cluster, and then create the Kibana Kubernetes Service and Deployment. To conclude, we'll set up Fluentd as a DaemonSet so it runs on every Kubernetes worker node.

Prerequisites

Before you begin with this guide, ensure you have the following available to you:

- A Kubernetes 1.10+ cluster with role-based access control (RBAC) enabled
 - Ensure your cluster has enough resources available to roll out the EFK stack, and if not scale your cluster by adding worker nodes. We'll be deploying a 3-Pod Elasticsearch cluster (you can scale this down to 1 if necessary), as well as a single Kibana Pod. Every worker node will also run a Fluentd Pod. The cluster in this guide consists of 3 worker nodes and a managed control plane.
- The `kubectl` command-line tool installed on your local machine, configured to connect to your cluster. You can read more about installing `kubectl` [in the official documentation](#).

Once you have these components set up, you're ready to begin with this guide.

Step 1 — Creating a Namespace

Before we roll out an Elasticsearch cluster, we'll first create a Namespace into which we'll install all of our logging instrumentation. Kubernetes lets you separate objects running in your cluster using a "virtual cluster" abstraction called Namespaces. In this guide, we'll create a `kube-logging` namespace into which we'll install the EFK stack components.

SCROLL TO TOP

This Namespace will also allow us to quickly clean up and remove the logging stack without any loss of function to the Kubernetes cluster.

To begin, first investigate the existing Namespaces in your cluster using `kubectl`:

```
$ kubectl get namespaces
```

You should see the following three initial Namespaces, which come preinstalled with your Kubernetes cluster:

Output

NAME	STATUS	AGE
default	Active	5m
kube-system	Active	5m
kube-public	Active	5m

The `default` Namespace houses objects that are created without specifying a Namespace. The `kube-system` Namespace contains objects created and used by the Kubernetes system, like `kube-dns`, `kube-proxy`, and `kubernetes-dashboard`. It's good practice to keep this Namespace clean and not pollute it with your application and instrumentation workloads.

The `kube-public` Namespace is another automatically created Namespace that can be used to store objects you'd like to be readable and accessible throughout the whole cluster, even to unauthenticated users.

To create the `kube-logging` Namespace, first open and edit a file called `kube-logging.yaml` using your favorite editor, such as `nano`:

```
$ nano kube-logging.yaml
```

Inside your editor, paste the following Namespace object YAML:

kube-logging.yaml

```
kind: Namespace
apiVersion: v1
metadata:
  name: kube-logging
```

SCROLL TO TOP

Then, save and close the file.

Here, we specify the Kubernetes object's `kind` as a `Namespace` object. To learn more about `Namespace` objects, consult the [Namespaces Walkthrough](#) in the official Kubernetes documentation. We also specify the Kubernetes API version used to create the object (`v1`), and give it a `name`, `kube-logging`.

Once you've created the `kube-logging.yaml` `Namespace` object file, create the `Namespace` using `kubectl create` with the `-f` filename flag:

```
$ kubectl create -f kube-logging.yaml
```

You should see the following output:

Output

```
namespace/kube-logging created
```

You can then confirm that the `Namespace` was successfully created:

```
$ kubectl get namespaces
```

At this point, you should see the new `kube-logging` `Namespace`:

Output

NAME	STATUS	AGE
default	Active	23m
kube-logging	Active	1m
kube-public	Active	23m
kube-system	Active	23m

We can now deploy an Elasticsearch cluster into this isolated logging `Namespace`.

Step 2 — Creating the Elasticsearch StatefulSet

Now that we've created a `Namespace` to house our logging stack, we can begin rolling out its various components. We'll first begin by deploying a 3-node Elasticsearch cluster.

In this guide, we use 3 Elasticsearch Nodes to avoid the "split-brain" issue that occurs in highly-available, multi-node clusters. [SCROLL TO TOP](#), split-brain" is what arises when one

or more nodes can't communicate with the others, and several "split" masters get elected. With 3 nodes, if one gets disconnected from the cluster temporarily, the other two nodes can elect a new master and the cluster can continue functioning while the last node attempts to rejoin. To learn more, consult [A new era for cluster coordination in Elasticsearch](#) and [Voting configurations](#).

Creating the Headless Service

To start, we'll create a headless Kubernetes service called `elasticsearch` that will define a DNS domain for the 3 Pods. A headless service does not perform load balancing or have a static IP; to learn more about headless services, consult the official [Kubernetes documentation](#).

Open a file called `elasticsearch_svc.yaml` using your favorite editor:

```
$ nano elasticsearch_svc.yaml
```

Paste in the following Kubernetes service YAML:

`elasticsearch_svc.yaml`

```
kind: Service
apiVersion: v1
metadata:
  name: elasticsearch
  namespace: kube-logging
  labels:
    app: elasticsearch
spec:
  selector:
    app: elasticsearch
  clusterIP: None
  ports:
    - port: 9200
      name: rest
    - port: 9300
      name: inter-node
```

Then, save and close the file.

We define a Service called `elasticsearch` in the `kube-logging` Namespace, and give it the `app: elasticsearch` label. We then set the `.spec.selector` to `app: elasticsearch` so that the Service finds all Pods with the `app: elasticsearch` label. When we associate our Elasticsearch StatefulSet with this Service, the Service will

[SCROLL TO TOP](#)

return DNS A records that point to Elasticsearch Pods with the `app: elasticsearch` label.

We then set `clusterIP: None`, which renders the service headless. Finally, we define ports `9200` and `9300` which are used to interact with the REST API, and for inter-node communication, respectively.

Create the service using `kubectl`:

```
$ kubectl create -f elasticsearch_svc.yaml
```

You should see the following output:

Output

```
service/elasticsearch created
```

Finally, double-check that the service was successfully created using `kubectl get`:

```
kubectl get services --namespace=kube-logging
```

You should see the following:

Output

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
elasticsearch	ClusterIP	None	<none>	9200/TCP,9300/TCP	26s

Now that we've set up our headless service and a stable `.elasticsearch.kube-logging.svc.cluster.local` domain for our Pods, we can go ahead and create the StatefulSet.

Creating the StatefulSet

A Kubernetes StatefulSet allows you to assign a stable identity to Pods and grant them stable, persistent storage. Elasticsearch requires stable storage to persist data across Pod rescheduling and restarts. To learn more about the StatefulSet workload, consult the [Statefulsets](#) page from the Kubernetes docs.

Open a file called `elasticsearch_s` [SCROLL TO TOP](#) in your favorite editor:

```
$ nano elasticsearch_statefulset.yaml
```

We will move through the StatefulSet object definition section by section, pasting blocks into this file.

Begin by pasting in the following block:

```
elasticsearch_statefulset.yaml
```

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: es-cluster
  namespace: kube-logging
spec:
  serviceName: elasticsearch
  replicas: 3
  selector:
    matchLabels:
      app: elasticsearch
  template:
    metadata:
      labels:
        app: elasticsearch
```

In this block, we define a StatefulSet called `es-cluster` in the `kube-logging` namespace. We then associate it with our previously created `elasticsearch` Service using the `serviceName` field. This ensures that each Pod in the StatefulSet will be accessible using the following DNS address:

`es-cluster-[0,1,2].elasticsearch.kube-logging.svc.cluster.local`, where `[0,1,2]` corresponds to the Pod's assigned integer ordinal.

We specify 3 `replicas` (Pods) and set the `matchLabels` selector to `app: elasticsearch`, which we then mirror in the `.spec.template.metadata` section. The `.spec.selector.matchLabels` and `.spec.template.metadata.labels` fields must match.

We can now move on to the object spec. Paste in the following block of YAML immediately below the preceding block:

```
elasticsearch_statefulset.yaml
```

SCROLL TO TOP

```
spec:
```

```

containers:
- name: elasticsearch
  image: docker.elastic.co/elasticsearch/elasticsearch:7.2.0
  resources:
    limits:
      cpu: 1000m
    requests:
      cpu: 100m
  ports:
  - containerPort: 9200
    name: rest
    protocol: TCP
  - containerPort: 9300
    name: inter-node
    protocol: TCP
  volumeMounts:
  - name: data
    mountPath: /usr/share/elasticsearch/data
  env:
  - name: cluster.name
    value: k8s-logs
  - name: node.name
    valueFrom:
      fieldRef:
        fieldPath: metadata.name
  - name: discovery.seed_hosts
    value: "es-cluster-0.elasticsearch,es-cluster-1.elasticsearch,es-cluster-2.elasticsearch"
  - name: cluster.initial_master_nodes
    value: "es-cluster-0,es-cluster-1,es-cluster-2"
  - name: ES_JAVA_OPTS
    value: "-Xms512m -Xmx512m"

```

Here we define the Pods in the StatefulSet. We name the containers `elasticsearch` and choose the `docker.elastic.co/elasticsearch/elasticsearch:7.2.0` Docker image. At this point, you may modify this image tag to correspond to your own internal Elasticsearch image, or a different version. Note that for the purposes of this guide, only Elasticsearch 7.2.0 has been tested.

We then use the `resources` field to specify that the container needs at least 0.1 vCPU guaranteed to it, and can burst up to 1 vCPU (which limits the Pod's resource usage when performing an initial large ingest or dealing with a load spike). You should modify these values depending on your anticipated load and available resources. To learn more about resource requests and limits, consult the official [Kubernetes Documentation](https://kubernetes.io/docs/concepts/configuration/resource-limits/).

We then open and name ports 9200 and 9300 for REST API and inter-node communication, respectively. We specify a `volumeMount` called `data` that will mount the PersistentVolume named `data` to the `/usr/share/elasticsearch/data` path

[SCROLL TO TOP](#)

`/usr/share/elasticsearch/data`. We will define the VolumeClaims for this StatefulSet in a later YAML block.

Finally, we set some environment variables in the container:

- `cluster.name`: The Elasticsearch cluster's name, which in this guide is `k8s-logs`.
- `node.name`: The node's name, which we set to the `.metadata.name` field using `valueFrom`. This will resolve to `es-cluster-[0,1,2]`, depending on the node's assigned ordinal.
- `discovery.seed_hosts`: This field sets a list of master-eligible nodes in the cluster that will seed the node discovery process. In this guide, thanks to the headless service we configured earlier, our Pods have domains of the form `es-cluster-[0,1,2].elasticsearch.kube-logging.svc.cluster.local`, so we set this variable accordingly. Using local namespace Kubernetes DNS resolution, we can shorten this to `es-cluster-[0,1,2].elasticsearch`. To learn more about Elasticsearch discovery, consult the official [Elasticsearch documentation](#).
- `cluster.initial_master_nodes`: This field also specifies a list of master-eligible nodes that will participate in the master election process. Note that for this field you should identify nodes by their `node.name`, and not their hostnames.
- `ES_JAVA_OPTS`: Here we set this to `-Xms512m -Xmx512m` which tells the JVM to use a minimum and maximum heap size of 512 MB. You should tune these parameters depending on your cluster's resource availability and needs. To learn more, consult [Setting the heap size](#).

The next block we'll paste in looks as follows:

elasticsearch_statefulset.yaml

```

. . .
  initContainers:
  - name: fix-permissions
    image: busybox
    command: ["sh", "-c", "chown -R 1000:1000 /usr/share/elasticsearch/data"]
    securityContext:
      privileged: true
    volumeMounts:
    - name: data
      mountPath: /usr/share/elasticsearch/data
  - name: increase-vm-max-map
    image: busybox
    command: ["sysctl", "-w", "vm.max_map_count=262144"]
    securityContext:
      privileged: true
  - name: increase-fd-ulimit

```

SCROLL TO TOP

```
image: busybox
command: ["sh", "-c", "ulimit -n 65536"]
securityContext:
  privileged: true
```

In this block, we define several Init Containers that run before the main `elasticsearch` app container. These Init Containers each run to completion in the order they are defined. To learn more about Init Containers, consult the official [Kubernetes Documentation](#).

The first, named `fix-permissions`, runs a `chown` command to change the owner and group of the Elasticsearch data directory to `1000:1000`, the Elasticsearch user's UID. By default Kubernetes mounts the data directory as `root`, which renders it inaccessible to Elasticsearch. To learn more about this step, consult Elasticsearch's ["Notes for production use and defaults."](#)

The second, named `increase-vm-max-map`, runs a command to increase the operating system's limits on `mmap` counts, which by default may be too low, resulting in out of memory errors. To learn more about this step, consult the official [Elasticsearch documentation](#).

The next Init Container to run is `increase-fd-ulimit`, which runs the `ulimit` command to increase the maximum number of open file descriptors. To learn more about this step, consult the ["Notes for Production Use and Defaults"](#) from the official Elasticsearch documentation.

Note: The Elasticsearch [Notes for Production Use](#) also mentions disabling swapping for performance reasons. Depending on your Kubernetes installation or provider, swapping may already be disabled. To check this, `exec` into a running container and run `cat /proc/swaps` to list active swap devices. If you see nothing there, swap is disabled.

Now that we've defined our main app container and the Init Containers that run before it to tune the container OS, we can add the final piece to our StatefulSet object definition file: the `volumeClaimTemplates`.

Paste in the following `volumeClaimTemplate` block:

`elasticsearch_statefulset.yaml`

```
...
volumeClaimTemplates:
  - metadata:
```

SCROLL TO TOP

```

    name: data
    labels:
      app: elasticsearch
spec:
  accessModes: [ "ReadWriteOnce" ]
  storageClassName: do-block-storage
  resources:
    requests:
      storage: 100Gi

```

In this block, we define the StatefulSet's `volumeClaimTemplates`. Kubernetes will use this to create PersistentVolumes for the Pods. In the block above, we name it `data` (which is the name we refer to in the `volumeMounts` defined previously), and give it the same `app: elasticsearch` label as our StatefulSet.

We then specify its access mode as `ReadWriteOnce`, which means that it can only be mounted as read-write by a single node. We define the storage class as `do-block-storage` in this guide since we use a DigitalOcean Kubernetes cluster for demonstration purposes. You should change this value depending on where you are running your Kubernetes cluster. To learn more, consult the [Persistent Volume](#) documentation.

Finally, we specify that we'd like each PersistentVolume to be 100GiB in size. You should adjust this value depending on your production needs.

The complete StatefulSet spec should look something like this:

elasticsearch_statefulset.yaml

```

apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: es-cluster
  namespace: kube-logging
spec:
  serviceName: elasticsearch
  replicas: 3
  selector:
    matchLabels:
      app: elasticsearch
  template:
    metadata:
      labels:
        app: elasticsearch
    spec:
      containers:
        - name: elasticsearch

```

SCROLL TO TOP

```

image: docker.elastic.co/elasticsearch/elasticsearch:7.2.0
resources:
  limits:
    cpu: 1000m
  requests:
    cpu: 100m
ports:
- containerPort: 9200
  name: rest
  protocol: TCP
- containerPort: 9300
  name: inter-node
  protocol: TCP
volumeMounts:
- name: data
  mountPath: /usr/share/elasticsearch/data
env:
- name: cluster.name
  value: k8s-logs
- name: node.name
  valueFrom:
    fieldRef:
      fieldPath: metadata.name
- name: discovery.seed_hosts
  value: "es-cluster-0.elasticsearch,es-cluster-1.elasticsearch,es-cluster-2.elasticsearch"
- name: cluster.initial_master_nodes
  value: "es-cluster-0,es-cluster-1,es-cluster-2"
- name: ES_JAVA_OPTS
  value: "-Xms512m -Xmx512m"
initContainers:
- name: fix-permissions
  image: busybox
  command: ["sh", "-c", "chown -R 1000:1000 /usr/share/elasticsearch/data"]
  securityContext:
    privileged: true
  volumeMounts:
  - name: data
    mountPath: /usr/share/elasticsearch/data
- name: increase-vm-max-map
  image: busybox
  command: ["sysctl", "-w", "vm.max_map_count=262144"]
  securityContext:
    privileged: true
- name: increase-fd-ulimit
  image: busybox
  command: ["sh", "-c", "ulimit -n 65536"]
  securityContext:
    privileged: true
volumeClaimTemplates:
- metadata:
  name: data
  labels:
    app: elasticsearch
spec:
  accessModes: [ "ReadWriteOnce" ]

```

[SCROLL TO TOP](#)

```
storageClassName: do-block-storage
resources:
  requests:
    storage: 100Gi
```

Once you're satisfied with your Elasticsearch configuration, save and close the file.

Now, deploy the StatefulSet using `kubectl`:

```
$ kubectl create -f elasticsearch_statefulset.yaml
```

You should see the following output:

Output

```
statefulset.apps/es-cluster created
```

You can monitor the StatefulSet as it is rolled out using `kubectl rollout status`:

```
$ kubectl rollout status sts/es-cluster --namespace=kube-logging
```

You should see the following output as the cluster is rolled out:

Output

```
Waiting for 3 pods to be ready...
Waiting for 2 pods to be ready...
Waiting for 1 pods to be ready...
partitioned roll out complete: 3 new pods have been updated...
```

Once all the Pods have been deployed, you can check that your Elasticsearch cluster is functioning correctly by performing a request against the REST API.

To do so, first forward the local port `9200` to the port `9200` on one of the Elasticsearch nodes (`es-cluster-0`) using `kubectl port-forward`:

```
$ kubectl port-forward es-cluster-0 9200:9200 --namespace=kube-logging
```

Then, in a separate terminal window, perform a `curl` request against the REST API:

SCROLL TO TOP

```
$ curl http://localhost:9200/_cluster/state?pretty
```

You should see the following output:

Output

```
{
  "cluster_name" : "k8s-logs",
  "compressed_size_in_bytes" : 348,
  "cluster_uuid" : "QD06dK7CQgids-GQZooNVw",
  "version" : 3,
  "state_uuid" : "mjNIWXAzQVuxNNOQ7xR-qg",
  "master_node" : "IdM5B7cUQWqFgIHXBp0JDg",
  "blocks" : { },
  "nodes" : {
    "u7DoTpMmSCix0oictzHItA" : {
      "name" : "es-cluster-1",
      "ephemeral_id" : "ZlBflnXKRC4RvEACHIVdg",
      "transport_address" : "10.244.8.2:9300",
      "attributes" : { }
    },
    "IdM5B7cUQWqFgIHXBp0JDg" : {
      "name" : "es-cluster-0",
      "ephemeral_id" : "JTk1FDdFQuWbSFAtBxdxAQ",
      "transport_address" : "10.244.44.3:9300",
      "attributes" : { }
    },
    "R8E7xcSUSbGbgrhAdyAKmQ" : {
      "name" : "es-cluster-2",
      "ephemeral_id" : "9wv6ke71Qqy9vk2LgJTqaA",
      "transport_address" : "10.244.40.4:9300",
      "attributes" : { }
    }
  },
  ...
}
```

This indicates that our Elasticsearch cluster `k8s-logs` has successfully been created with 3 nodes: `es-cluster-0`, `es-cluster-1`, and `es-cluster-2`. The current master node is `es-cluster-0`.

Now that your Elasticsearch cluster is up and running, you can move on to setting up a Kibana frontend for it.

Step 3 — Creating the Kibana SCROLL TO TOP Service

To launch Kibana on Kubernetes, we'll create a Service called `kibana`, and a Deployment consisting of one Pod replica. You can scale the number of replicas depending on your production needs, and optionally specify a `LoadBalancer` type for the Service to load balance requests across the Deployment pods.

This time, we'll create the Service and Deployment in the same file. Open up a file called `kibana.yaml` in your favorite editor:

```
$ nano kibana.yaml
```

Paste in the following service spec:

kibana.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: kibana
  namespace: kube-logging
  labels:
    app: kibana
spec:
  ports:
    - port: 5601
  selector:
    app: kibana
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kibana
  namespace: kube-logging
  labels:
    app: kibana
spec:
  replicas: 1
  selector:
    matchLabels:
      app: kibana
  template:
    metadata:
      labels:
        app: kibana
    spec:
      containers:
        - name: kibana
          image: docker.elastic.co/kibana/kibana:7.2.0
          resources:
            limits:
```

SCROLL TO TOP

```
    cpu: 1000m
  requests:
    cpu: 100m
  env:
    - name: ELASTICSEARCH_URL
      value: http://elasticsearch:9200
  ports:
    - containerPort: 5601
```

Then, save and close the file.

In this spec we've defined a service called `kibana` in the `kube-logging` namespace, and gave it the `app: kibana` label.

We've also specified that it should be accessible on port `5601` and use the `app: kibana` label to select the Service's target Pods.

In the `Deployment` spec, we define a `Deployment` called `kibana` and specify that we'd like 1 Pod replica.

We use the `docker.elastic.co/kibana/kibana:7.2.0` image. At this point you may substitute your own private or public Kibana image to use.

We specify that we'd like at the very least 0.1 vCPU guaranteed to the Pod, bursting up to a limit of 1 vCPU. You may change these parameters depending on your anticipated load and available resources.

Next, we use the `ELASTICSEARCH_URL` environment variable to set the endpoint and port for the Elasticsearch cluster. Using Kubernetes DNS, this endpoint corresponds to its Service name `elasticsearch`. This domain will resolve to a list of IP addresses for the 3 Elasticsearch Pods. To learn more about Kubernetes DNS, consult [DNS for Services and Pods](#).

Finally, we set Kibana's container port to `5601`, to which the `kibana` Service will forward requests.

Once you're satisfied with your Kibana configuration, you can roll out the Service and Deployment using `kubectl`:

```
$ kubectl create -f kibana.yaml
```

SCROLL TO TOP

You should see the following output:

Output

```
service/kibana created
deployment.apps/kibana created
```

You can check that the rollout succeeded by running the following command:

```
$ kubectl rollout status deployment/kibana --namespace=kube-logging
```

You should see the following output:

Output

```
deployment "kibana" successfully rolled out
```

To access the Kibana interface, we'll once again forward a local port to the Kubernetes node running Kibana. Grab the Kibana Pod details using `kubectl get`:

```
$ kubectl get pods --namespace=kube-logging
```

Output

NAME	READY	STATUS	RESTARTS	AGE
es-cluster-0	1/1	Running	0	55m
es-cluster-1	1/1	Running	0	54m
es-cluster-2	1/1	Running	0	54m
kibana-6c9fb4b5b7-plbg2	1/1	Running	0	4m27s

Here we observe that our Kibana Pod is called `kibana-6c9fb4b5b7-plbg2`.

Forward the local port `5601` to port `5601` on this Pod:

```
$ kubectl port-forward kibana-6c9fb4b5b7-plbg2 5601:5601 --namespace=kube-logging
```

You should see the following output:

Output

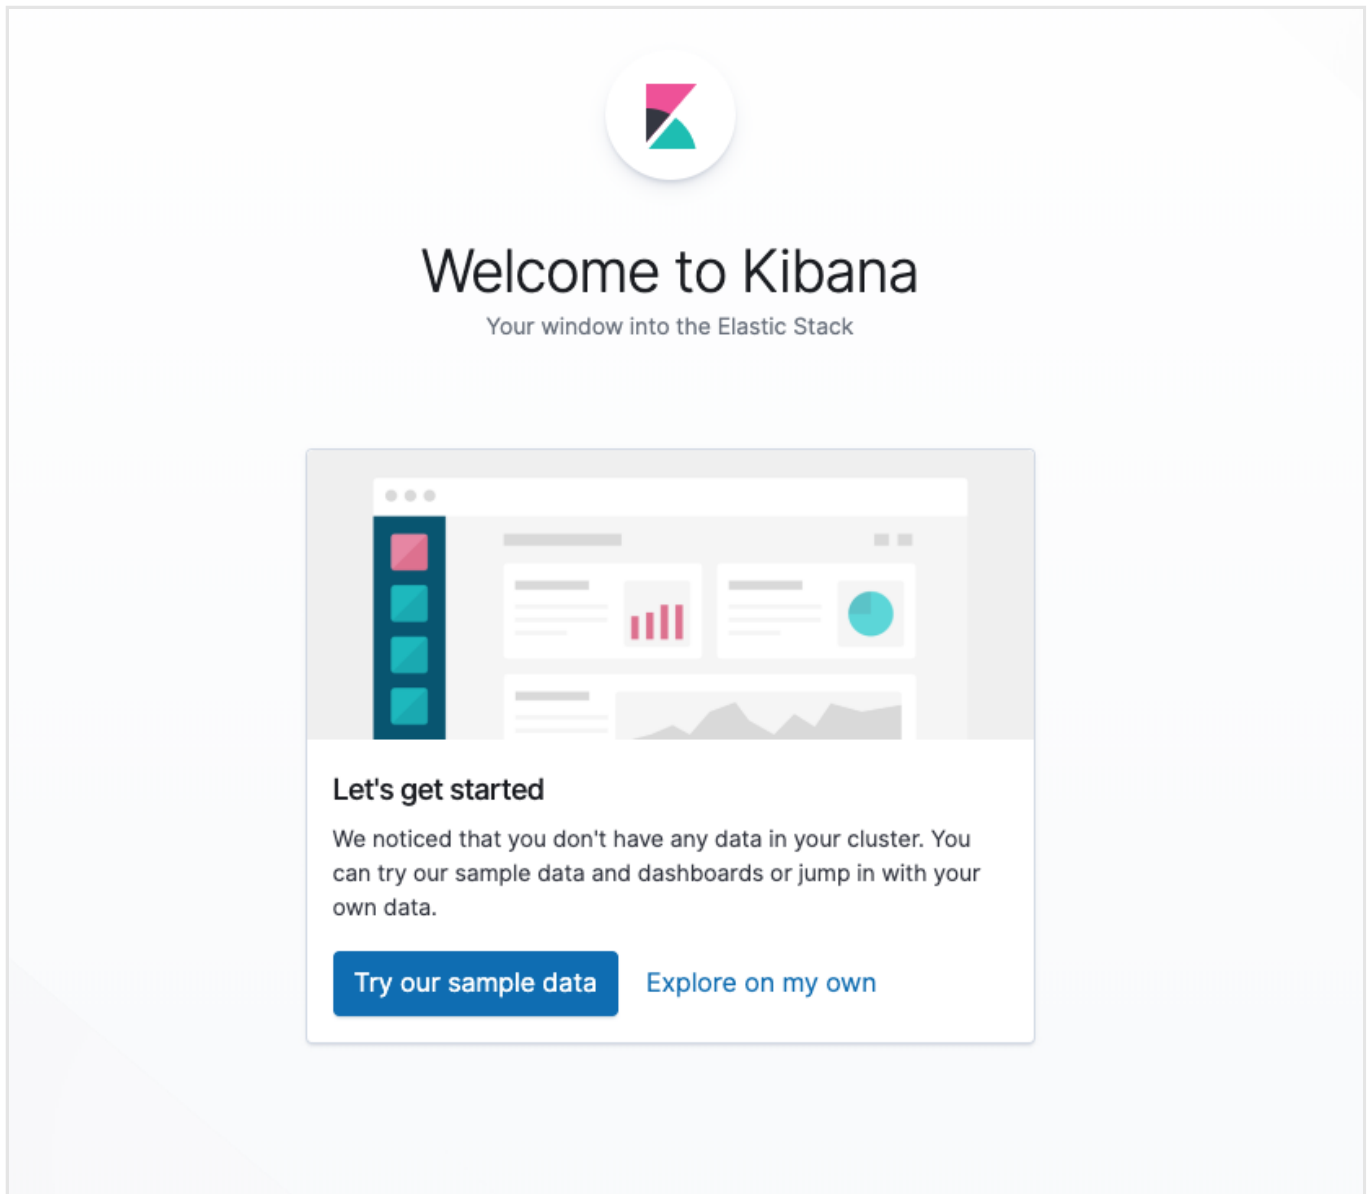
```
Forwarding from 127.0.0.1:5601 -> 5601
Forwarding from [::1]:5601 -> 5601
```

SCROLL TO TOP

Now, in your web browser, visit the following URL:

`http://localhost:5601`

If you see the following Kibana welcome page, you've successfully deployed Kibana into your Kubernetes cluster:



You can now move on to rolling out the final component of the EFK stack: the log collector, Fluentd.

Step 4 — Creating the Fluentd DaemonSet

In this guide, we'll set up Fluentd as a DaemonSet, which is a Kubernetes workload type that runs a copy of a given Pod on each Node in the Kubernetes cluster. Using this DaemonSet controller, we'll roll out a [Fluentd agent Pod](#) on every node in our

[SCROLL TO TOP](#)

cluster. To learn more about this logging architecture, consult [“Using a node logging agent”](#) from the official Kubernetes docs.

In Kubernetes, containerized applications that log to `stdout` and `stderr` have their log streams captured and redirected to JSON files on the nodes. The Fluentd Pod will tail these log files, filter log events, transform the log data, and ship it off to the Elasticsearch logging backend we deployed in [Step 2](#).

In addition to container logs, the Fluentd agent will tail Kubernetes system component logs like kubelet, kube-proxy, and Docker logs. To see a full list of sources tailed by the Fluentd logging agent, consult the [kubernetes.conf](#) file used to configure the logging agent. To learn more about logging in Kubernetes clusters, consult [“Logging at the node level”](#) from the official Kubernetes documentation.

Begin by opening a file called `fluentd.yaml` in your favorite text editor:

```
$ nano fluentd.yaml
```

Once again, we'll paste in the Kubernetes object definitions block by block, providing context as we go along. In this guide, we use the [Fluentd DaemonSet spec](#) provided by the Fluentd maintainers. Another helpful resource provided by the Fluentd maintainers is [Kuberentes Fluentd](#).

First, paste in the following ServiceAccount definition:

fluentd.yaml

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: fluentd
  namespace: kube-logging
  labels:
    app: fluentd
```

Here, we create a Service Account called `fluentd` that the Fluentd Pods will use to access the Kubernetes API. We create it in the `kube-logging` Namespace and once again give it the label `app: fluentd`. To learn more about Service Accounts in Kubernetes, consult [Configure Service Accounts for Pods](#) in the official Kubernetes docs.

Next, paste in the following Cluster [SCROLL TO TOP](#)

fluentd.yaml

```

. . .
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: fluentd
  labels:
    app: fluentd
rules:
- apiGroups:
  - ""
  resources:
  - pods
  - namespaces
  verbs:
  - get
  - list
  - watch

```

Here we define a ClusterRole called `fluentd` to which we grant the `get`, `list`, and `watch` permissions on the `pods` and `namespaces` objects. ClusterRoles allow you to grant access to cluster-scoped Kubernetes resources like Nodes. To learn more about Role-Based Access Control and Cluster Roles, consult [Using RBAC Authorization](#) from the official Kubernetes documentation.

Now, paste in the following ClusterRoleBinding block:

fluentd.yaml

```

. . .
---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: fluentd
roleRef:
  kind: ClusterRole
  name: fluentd
  apiGroup: rbac.authorization.k8s.io
subjects:
- kind: ServiceAccount
  name: fluentd
  namespace: kube-logging

```

In this block, we define a ClusterRoleBinding called `fluentd` which binds the `fluentd` ClusterRole to the `fluentd` ServiceAccount. This grants the `fluentd`

SCROLL TO TOP

ServiceAccount the permissions listed in the `fluentd` Cluster Role.

At this point we can begin pasting in the actual DaemonSet spec:

fluentd.yaml

```

. . .
---
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd
  namespace: kube-logging
  labels:
    app: fluentd

```

Here, we define a DaemonSet called `fluentd` in the `kube-logging` Namespace and give it the `app: fluentd` label.

Next, paste in the following section:

fluentd.yaml

```

. . .
spec:
  selector:
    matchLabels:
      app: fluentd
  template:
    metadata:
      labels:
        app: fluentd
    spec:
      serviceAccount: fluentd
      serviceAccountName: fluentd
      tolerations:
        - key: node-role.kubernetes.io/master
          effect: NoSchedule
      containers:
        - name: fluentd
          image: fluent/fluentd-kubernetes-daemonset:v1.4.2-debian-elasticsearch-1.1
          env:
            - name: FLUENT_ELASTICSEARCH_HOST
              value: "elasticsearch.kube-logging.svc.cluster.local"
            - name: FLUENT_ELASTICSEARCH_PORT
              value: "9200"
            - name: FLUENT_ELASTICSEARCH_SCHEME
              value: "http"
            - name: FLUENTD_SYSTEMD_ENABLED
              value: disable

```

Here, we match the `app: fluentd` label defined in `.metadata.labels` and then assign the DaemonSet the `fluentd` Service Account. We also select the `app: fluentd` as the Pods managed by this DaemonSet.

Next, we define a `NoSchedule` toleration to match the equivalent taint on Kubernetes master nodes. This will ensure that the DaemonSet also gets rolled out to the Kubernetes masters. If you don't want to run a Fluentd Pod on your master nodes, remove this toleration. To learn more about Kubernetes taints and tolerations, consult ["Taints and Tolerations"](#) from the official Kubernetes docs.

Next, we begin defining the Pod container, which we call `fluentd`.

We use the [official v1.4.2 Debian image](#) provided by the Fluentd maintainers. If you'd like to use your own private or public Fluentd image, or use a different image version, modify the `image` tag in the container spec. The Dockerfile and contents of this image are available in Fluentd's [fluentd-kubernetes-daemonset Github repo](#).

Next, we configure Fluentd using some environment variables:

- `FLUENT_ELASTICSEARCH_HOST`: We set this to the Elasticsearch headless Service address defined earlier: `elasticsearch.kube-logging.svc.cluster.local`. This will resolve to a list of IP addresses for the 3 Elasticsearch Pods. The actual Elasticsearch host will most likely be the first IP address returned in this list. To distribute logs across the cluster, you will need to modify the configuration for Fluentd's Elasticsearch Output plugin. To learn more about this plugin, consult [Elasticsearch Output Plugin](#).
- `FLUENT_ELASTICSEARCH_PORT`: We set this to the Elasticsearch port we configured earlier, `9200`.
- `FLUENT_ELASTICSEARCH_SCHEME`: We set this to `http`.
- `FLUENTD_SYSTEMD_CONF`: We set this to `disable` to suppress output related to `systemd` not being set up in the container.

Finally, paste in the following section:

fluentd.yaml

...

```
resources:
  limits:
    memory: 512Mi
  requests:
```

SCROLL TO TOP

```

    cpu: 100m
    memory: 200Mi
  volumeMounts:
  - name: varlog
    mountPath: /var/log
  - name: varlibdockercontainers
    mountPath: /var/lib/docker/containers
    readOnly: true
  terminationGracePeriodSeconds: 30
  volumes:
  - name: varlog
    hostPath:
      path: /var/log
  - name: varlibdockercontainers
    hostPath:
      path: /var/lib/docker/containers

```

Here we specify a 512 MiB memory limit on the FluentD Pod, and guarantee it 0.1vCPU and 200MiB of memory. You can tune these resource limits and requests depending on your anticipated log volume and available resources.

Next, we mount the `/var/log` and `/var/lib/docker/containers` host paths into the container using the `varlog` and `varlibdockercontainers` `volumeMounts`. These volumes are defined at the end of the block.

The final parameter we define in this block is `terminationGracePeriodSeconds`, which gives Fluentd 30 seconds to shut down gracefully upon receiving a `SIGTERM` signal. After 30 seconds, the containers are sent a `SIGKILL` signal. The default value for `terminationGracePeriodSeconds` is 30s, so in most cases this parameter can be omitted. To learn more about gracefully terminating Kubernetes workloads, consult Google's ["Kubernetes best practices: terminating with grace."](#)

The entire Fluentd spec should look something like this:

fluentd.yaml

```

apiVersion: v1
kind: ServiceAccount
metadata:
  name: fluentd
  namespace: kube-logging
  labels:
    app: fluentd
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:

```

SCROLL TO TOP

```

    name: fluentd
    labels:
      app: fluentd
rules:
- apiGroups:
  - ""
  resources:
  - pods
  - namespaces
  verbs:
  - get
  - list
  - watch
---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: fluentd
roleRef:
  kind: ClusterRole
  name: fluentd
  apiGroup: rbac.authorization.k8s.io
subjects:
- kind: ServiceAccount
  name: fluentd
  namespace: kube-logging
---
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd
  namespace: kube-logging
  labels:
    app: fluentd
spec:
  selector:
    matchLabels:
      app: fluentd
  template:
    metadata:
      labels:
        app: fluentd
    spec:
      serviceAccount: fluentd
      serviceAccountName: fluentd
      tolerations:
      - key: node-role.kubernetes.io/master
        effect: NoSchedule
      containers:
      - name: fluentd
        image: fluent/fluentd-kubernetes-daemonset:v1.4.2-debian-elasticsearch-1.0
        env:
          - name: FLUENT_ELASTICSEARCH_HOST
            value: "elasticsearch-logging.kube-logging.svc.cluster.local"
          - name: FLUENT_ELASTICSEARCH_PORT

```



```

    value: "9200"
  - name: FLUENT_ELASTICSEARCH_SCHEME
    value: "http"
  - name: FLUENTD_SYSTEMD_CONF
    value: disable
resources:
  limits:
    memory: 512Mi
  requests:
    cpu: 100m
    memory: 200Mi
volumeMounts:
  - name: varlog
    mountPath: /var/log
  - name: varlibdockercontainers
    mountPath: /var/lib/docker/containers
    readOnly: true
terminationGracePeriodSeconds: 30
volumes:
  - name: varlog
    hostPath:
      path: /var/log
  - name: varlibdockercontainers
    hostPath:
      path: /var/lib/docker/containers

```

Once you've finished configuring the Fluentd DaemonSet, save and close the file.

Now, roll out the DaemonSet using `kubectl`:

```
$ kubectl create -f fluentd.yaml
```

You should see the following output:

Output

```

serviceaccount/fluentd created
clusterrole.rbac.authorization.k8s.io/fluentd created
clusterrolebinding.rbac.authorization.k8s.io/fluentd created
daemonset.extensions/fluentd created

```

Verify that your DaemonSet rolled out successfully using `kubectl`:

```
$ kubectl get ds --namespace=kube-logging
```

You should see the following status [SCROLL TO TOP](#)

Output

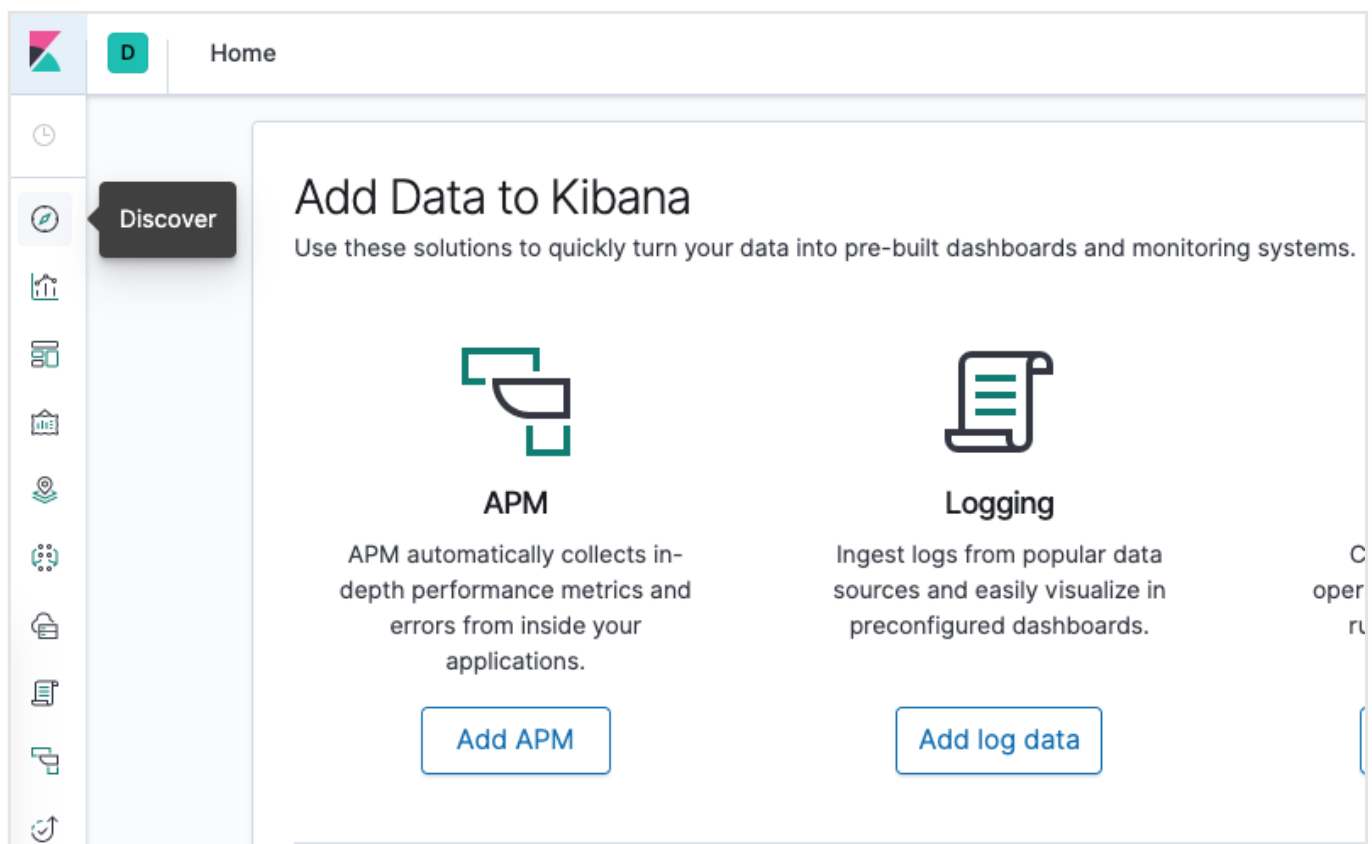
NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
fluentd	3	3	3	3	3	<none>	58s

This indicates that there are 3 `fluentd` Pods running, which corresponds to the number of nodes in our Kubernetes cluster.

We can now check Kibana to verify that log data is being properly collected and shipped to Elasticsearch.

With the `kubectl port-forward` still open, navigate to `http://localhost:5601`.

Click on **Discover** in the left-hand navigation menu:



You should see the following configuration window:

Management / Index patterns / Create index pattern

In order to visualize and explore data in Kibana, you'll need to create an index pattern to retrieve data from Elasticsearch.

Create index pattern

Kibana uses index patterns to retrieve data from Elasticsearch indices for things like visualizations.

☐ Include system indices

Step 1 of 2: Define index pattern

Index pattern

index-name-*

You can use a * as a wildcard in your index pattern. You can't use spaces or the characters \, /, ?, ", <, >, |.

No Elasticsearch indices match your pattern.

logstash-2019.07.08

Rows per page: 10

[Next step](#)

This allows you to define the Elasticsearch indices you'd like to explore in Kibana. To learn more, consult [Defining your index patterns](#) in the official Kibana docs. For now, we'll just use the `logstash-*` wildcard pattern to capture all the log data in our Elasticsearch cluster. Enter `logstash-*` in the text box and click on **Next step**.

You'll then be brought to the following page:

Management / Index patterns / Create index pattern

Create index pattern

Kibana uses index patterns to retrieve data from Elasticsearch indices for things like visualizations.

☐ Include system indices

Step 2 of 2: Configure settings

You've defined `logstash-*` as your index pattern. Now you can specify some settings before we create it.

Time Filter field name Refresh

The Time Filter will use this field to filter your data by time. You can choose not to have a time field, but you will not be able to narrow down your data by a time range.

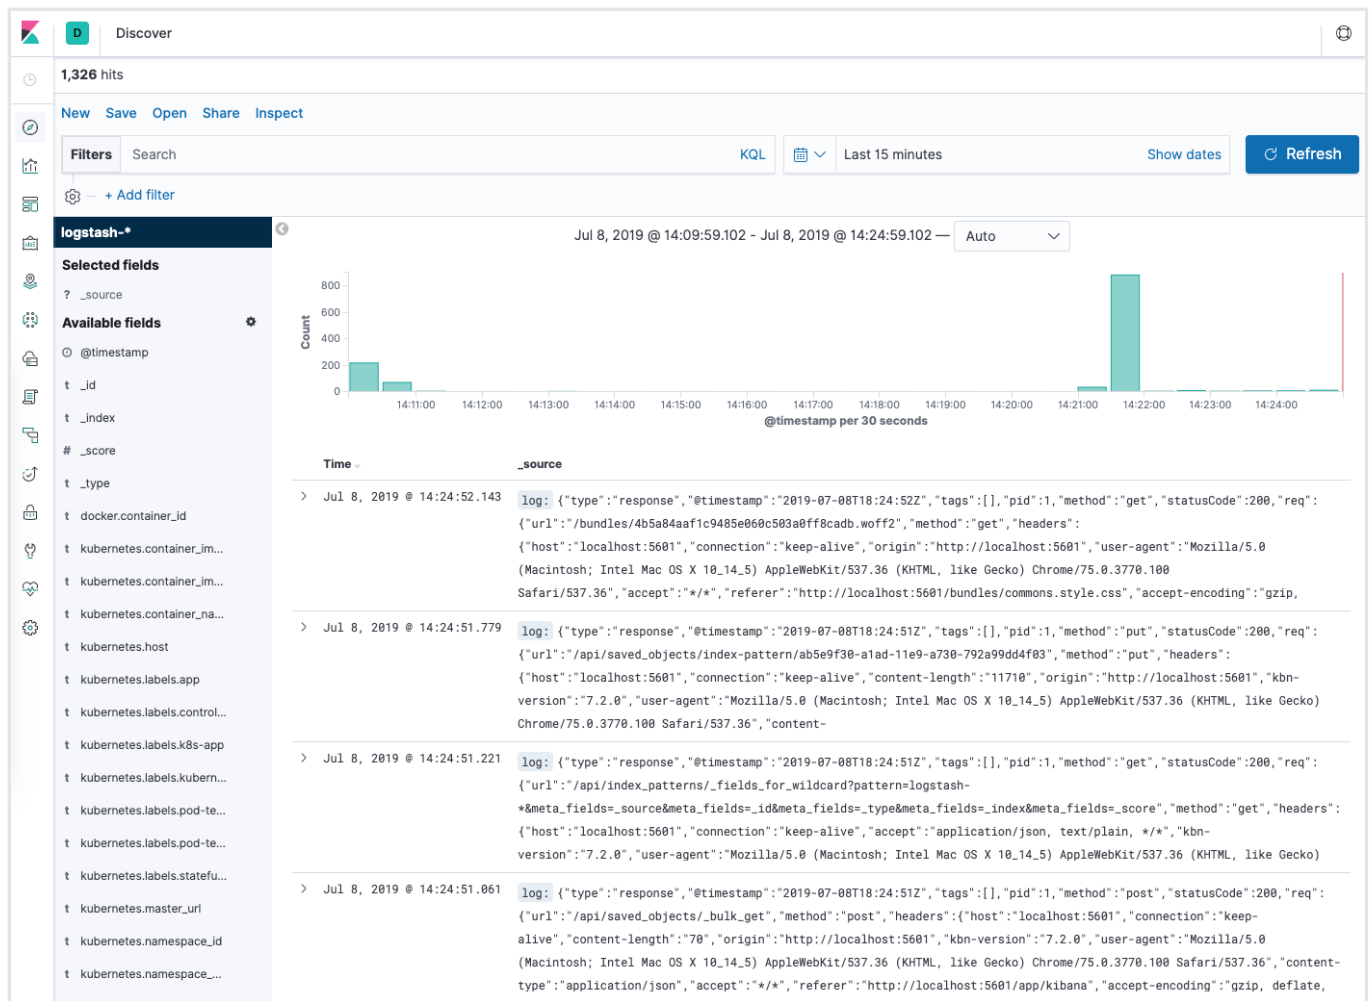
[Show advanced options](#)

[Back](#) [Create index pattern](#)

This allows you to configure which field Kibana will use to filter log data by time. In the dropdown, select the `@timestamp` field, and hit **Create index pattern**.

Now, hit **Discover** in the left hand navigation bar. [SCROLL TO TOP](#)

You should see a histogram graph and some recent log entries:



At this point you've successfully configured and rolled out the EFK stack on your Kubernetes cluster. To learn how to use Kibana to analyze your log data, consult the [Kibana User Guide](#).

In the next optional section, we'll deploy a simple counter Pod that prints numbers to stdout, and find its logs in Kibana.

Step 5 (Optional) — Testing Container Logging

To demonstrate a basic Kibana use case of exploring the latest logs for a given Pod, we'll deploy a minimal counter Pod that prints sequential numbers to stdout.

Let's begin by creating the Pod. Open up a file called `counter.yaml` in your favorite editor:

```
$ nano counter.yaml
```

SCROLL TO TOP

Then, paste in the following Pod spec.

counter.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: counter
spec:
  containers:
  - name: count
    image: busybox
    args: [/bin/sh, -c,
          'i=0; while true; do echo "$i: $(date)"; i=$((i+1)); sleep 1; done']
```

Save and close the file.

This is a minimal Pod called `counter` that runs a `while` loop, printing numbers sequentially.

Deploy the `counter` Pod using `kubectl`:

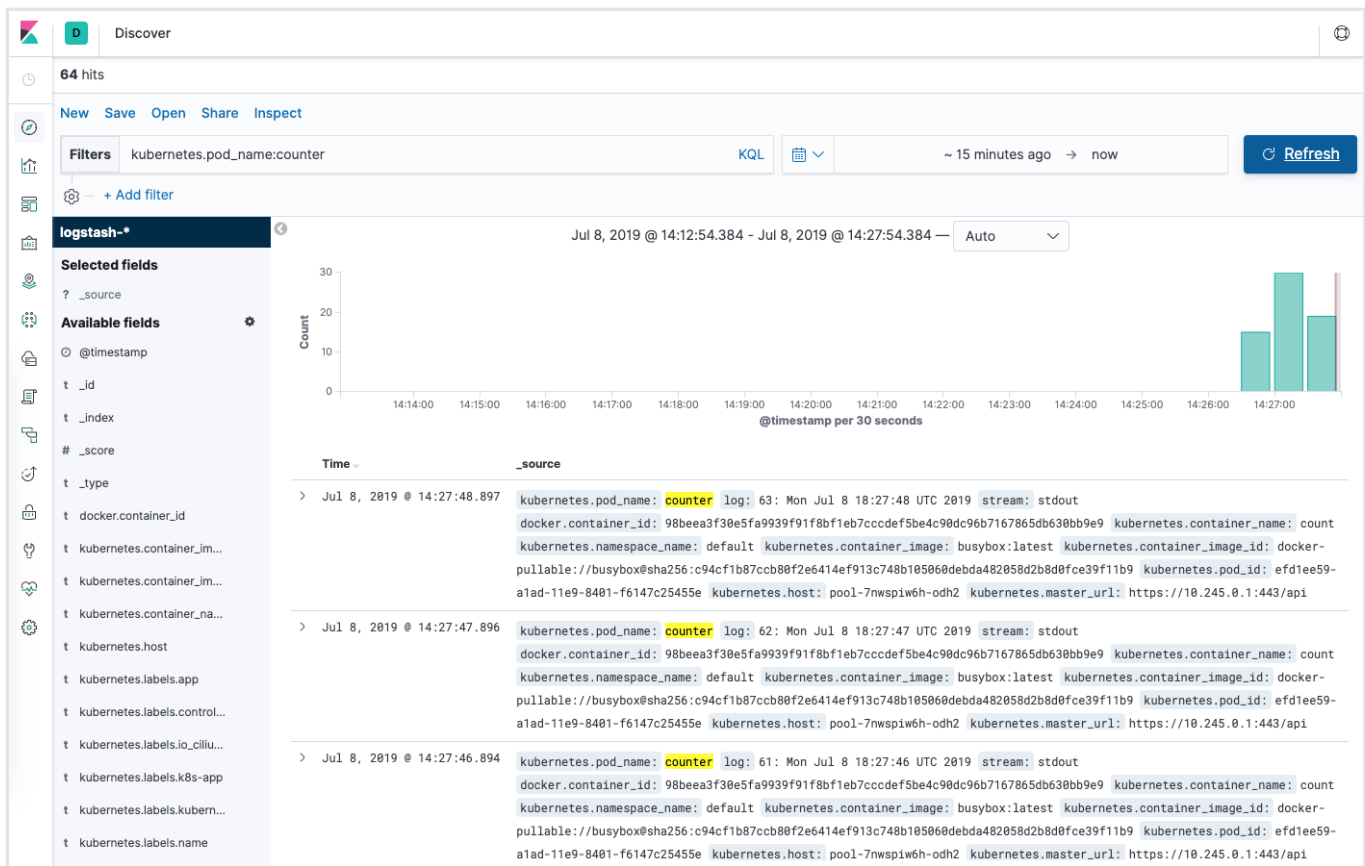
```
$ kubectl create -f counter.yaml
```

Once the Pod has been created and is running, navigate back to your Kibana dashboard.

From the **Discover** page, in the search bar enter `kubernetes.pod_name:counter`. This filters the log data for Pods named `counter`.

You should then see a list of log entries for the `counter` Pod:

SCROLL TO TOP



You can click into any of the log entries to see additional metadata like the container name, Kubernetes node, Namespace, and more.

Conclusion

In this guide we've demonstrated how to set up and configure Elasticsearch, Fluentd, and Kibana on a Kubernetes cluster. We've used a minimal logging architecture that consists of a single logging agent Pod running on each Kubernetes worker node.

Before deploying this logging stack into your production Kubernetes cluster, it's best to tune the resource requirements and limits as indicated throughout this guide. You may also want to set up [X-Pack](#) to enable built-in monitoring and security features.

The logging architecture we've used here consists of 3 Elasticsearch Pods, a single Kibana Pod (not load-balanced), and a set of Fluentd Pods rolled out as a DaemonSet. You may wish to scale this setup depending on your production use case. To learn more about scaling your Elasticsearch and Kibana stack, consult [Scaling Elasticsearch](#).

Kubernetes also allows for more complex logging agent architectures that may better suit your use case. To learn more, consult [Logging Architecture](#) from the Kubernetes docs.

Was this helpful?

Yes

No

SCROLL TO TOP



[Report an issue](#)

About the authors



Hanif Jetha

DevOps Engineer, Technical Writer
and Editor

Still looking for an answer?

[Ask a question](#)[Search for more help](#)

RELATED



DigitalOcean Spaces: Simple Object Storage
[Product](#)

How DigitalOcean Manages, Monitors, and Secures Kubernetes for App Platform

[Tutorial](#)

How To Install Apache Kafka on Ubuntu 20.04

[Tutorial](#)[SCROLL TO TOP](#)

54 Comments

Leave a comment...

Sign In to Comment

^ [ciaran4d51781530ee70807025](#) December 4, 2018



3 Very detailed and helpful tutorial thanks.

[Reply](#) [Report](#)

^ [ciaran4d51781530ee70807025](#) December 4, 2018



4 However, I'm having an issue on Digital Ocean's Kubernetes Service, I've tried on a new cluster, v1.12.3 and each time the elastic search stateful set fails due to issues with pvc's,

e.g. Warning FailedScheduling 16m default-scheduler pod has unbound immediate PersistentVolumeClaims

[Reply](#) [Report](#)

^ [leianivey](#) December 13, 2018



2 same issue.

[Reply](#) [Report](#)

^ [hjet](#)  December 19, 2018



2 Hi and thanks for your feedback!!


I just tested this on a DigitalOcean Kubernetes cluster (1.12.3) and was not able to replicate your issue (the ES StatefulSet was successfully created and I was able to curl the API) .

Does this occur intermittently for you? How large is your cluster, and what type of Nodes are you using? How much memory and vCPU do those Nodes have? Do you have multiple Node pools?

SCROLL TO TOP

Thanks!


[Reply](#) [Report](#)

 [ciaran4d51781530ee70807025](#) December 23, 2018
2 Hi,

I just retried on 1.12.3 with 3 nodes with 2GB RAM each and it worked. I don't remember my previous configuration but I suspect I had 1 node pool with 2 nodes with 1GB RAM each and 1 node pool with a single node with 4GB RAM.

Thanks for following up.

[Reply](#) [Report](#)

 [radugabor](#) October 26, 2019
0 You need to create StorgeClass and PersistentVolume:


```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: data
  namespace: kube-logging
  provisioner: kubernetes.io/no-provisioner
  volumeBindingMode: WaitForFirstConsumer
```

Create path /mnt/disk/vol1 with CHMOD 777

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: data
  namespace: kube-logging
spec:
  capacity:
    storage: 5Gi
  accessModes:
```

- ReadWriteOnce persistentVolumeReclaimPolicy: Retain storageClassName: data local: path: /mnt/disk/vol1 nodeAffinity: required: nodeSelectorTerms:
 - matchExpressions:
 - key: kubernetes.io/hostname operator: In values:
 - your-node-name

[Reply](#) [Report](#)

 [BuddyCasino](#) February 28, 2019
1

[SCROLL TO TOP](#)

This is the only detailed, correct and up-to-date tutorial on setting up a K8S Elasticsearch cluster on whole internet. Thanks!

[Reply](#) [Report](#)

^ [poojajagdale18](#) March 14, 2019

1 `curl http://localhost:9200/_cluster/state?pretty` gives me output as:

```
"error" : {
  "rootcause" : [
    {
      "type" : "mastemotdiscoveredexception",
      "reason" : null
    }
  ],
  "type" : "masternotdiscovered_exception",
  "reason" : null
},
"status" : 503
```

[Reply](#) [Report](#)

^ [hbinduni](#) May 19, 2019

0 it seem you forgot to deploy the service:

```
kubectl create -f elasticsearch_svc.yaml
```

[Reply](#) [Report](#)

^ [mishaveldhoen](#) May 28, 2019

0 This exception also occurs when you try to deploy ES v7 instead of v6 with the current settings. If this is the case, then you need to add the environment variable `cluster.initial_master_nodes`.

[Reply](#) [Report](#)

^ [rohitdharmapuri](#) May 18, 2020

0 My EFK stack is up and while running the below port forwards
`kubectl port-forward kibana-df8d558dd-27j9m 5601:5601 --namespace=logging`
`kubectl port-forward es-cluster-0 9200:9200 --namespace=logging`

I'm accessing my kibana on my chrome at this <http://localhost:5601/>

I'm getting this error intermittently `{"statusCode":503,"error":"Service Unavailable","message":"bl` [UNAVAILABLE/1/state not recovered /](#)

[SCROLL TO TOP](#)

initialized];: [clusterblockexception] blocked by: [SERVICEUNAVAILABLE/1/state not recovered / initialized];"} }

I made sure that I have the Kibana service up and ES has the following configuration

```
containers:
- env:
  - name: cluster.name
    value: k8s-logs
  - name: node.name
    valueFrom:
      fieldRef:
        apiVersion: v1
        fieldPath: metadata.name
  - name: discovery.seed_hosts
    value: es-cluster-0.elasticsearch,es-cluster-1.elasticsearch,es-
  - name: cluster.initial_master_nodes
    value: es-cluster-0,es-cluster-1,es-cluster-2
  - name: ES_JAVA_OPTS
    value: -Xms512m -Xmx512m
  image: gcr.io/****-****/efk/elasticsearch/elasticsearch:7.7.0
  imagePullPolicy: IfNotPresent
  name: elasticsearch
  ports:
  - containerPort: 9200
    name: rest
    protocol: TCP
  - containerPort: 9300
    name: inter-node
    protocol: TCP````
```

[Reply](#) [Report](#)

^ [jjohnson53b4c0bcd9f28deb44](#) April 24, 2019

0 One can also consider Fluent-bit. The bit variance has lighter functionality than Fluentd and for many its smaller size may make more sense.

[A Katacoda scenario is also available to walk you through setting up EFK here.](#)

[Reply](#) [Report](#)

^ [johnhalbert](#) May 21, 2019

0 Very detailed, thank you. Even with the recommended settings for ES_JAVA_OPTS , my Elasticsearch containers are dying due to OOM error. Looking at the available resources on the nodes where the pod is schedu [SCROLL TO TOP](#) e an issue. Any advice?

[Reply](#) [Report](#) [psuchy2](#) May 29, 2019 0 Hello, try to increase limit on your container resource. I had the same issue:

```
resources:
  limits:
    cpu: 1000m
    memory: 2.5G
```

[Reply](#) [Report](#) [mishaveldhoen](#) May 28, 2019 0 Great tutorial, thanks!


In case you're trying to deploy elasticsearch v7.x.x using these settings, no master node will be found. ES v7 introduces the `cluster.initial_master_nodes` setting, which will need to correspond exactly to the `node.name`, see [the ES reference](#), e.g., the node name must be `es-cluster-#.elasticsearch`.

One way to fix it is by changing

```
- name: node.name
  valueFrom:
    fieldRef:
      fieldPath: metadata.name
```

into

```
- name: NODE_NAME
  valueFrom:
    fieldRef:
      fieldPath: metadata.name
- name: node.name
  value: "${NODE_NAME}.elasticsearch"
- name: cluster.initial_master_nodes
  value: "es-cluster-0.elasticsearch,es-cluster-1.elasticsearch,es-cluster-2.el"
```

[Reply](#) [Report](#) [nsaitov](#) July 9, 2019 2 Hi [@mishaveldhoen](#), i am trying to create an elasticsearch_statefulset, but it got stuck on[SCROLL TO TOP](#)

Waiting for 3 pods to be ready...

I follow almost the same yaml file, except part for env as u mentioned:

env:

- name: cluster.name

value: k8s-logs

- name: NODENAME

valueFrom:

fieldRef:

fieldPath: metadata.name

- name: node.name

value: "\$ (NODENAME).elasticsearch"

- name: cluster.initialmastermodes

value: "es-cluster-0.elasticsearch,es-cluster-1.elasticsearch,es-cluster-2.elasticsearch"

- name: ESJAVA_OPTS

value: "-Xms512m -Xmx512m"

Please, let me know if i am doing anything wrong and how to proceed... Thanks in advance!

[Reply](#) [Report](#)



[josesantos](#) October 24, 2019

2

Me too.

Have you solved?

[Reply](#) [Report](#)



[psuchy2](#) May 29, 2019

0 Hello, I followed this tutorial, but after a few days I started to receive following message and no new logs are send to elasticsearch.

```
2019-05-29 12:39:58 +0000 [warn]: emit transaction failed: error_class=Fluent
```

Any tips, how to properly tweak fluentd?

Thanks

[Reply](#) [Report](#)



[jeanandrewfuentes](#) June 25, 2019

1

Hi great details on the tutorial, however I'm still getting an error.

```
plugin:elasticsearch@6.4.3 Service Unavailable
```

[SCROLL TO TOP](#)

at the kibana part it says status red too. and upon check of the pods of elastic this is the error.

fatal error on the network layer

```
at org.elasticsearch.transport.netty4.Netty4Utils.maybeDie(Netty4Utils.j
at org.elasticsearch.transport.netty4.Netty4MessageChannelHandler.excepti
at io.netty.channel.AbstractChannelHandlerContext.invokeExceptionCaught(/
at io.netty.channel.AbstractChannelHandlerContext.invokeExceptionCaught(/
```

Also this is the logs when I tried to proxy and curl elastic search

```
kubectrl port-forward es-cluster-0 9200:9200 --namespace=kube-logging
Forwarding from 127.0.0.1:9200 -> 9200
Forwarding from [::1]:9200 -> 9200
Handling connection for 9200
E0625 10:11:26.617570    7659 portforward.go:400] an error occurred forwardir
```

[Reply](#) [Report](#)

^ [mostafashr](#) July 6, 2019

1 I have the same error "FailedScheduling 16m default-scheduler pod has unbound immediate".

I test the system and find-out if I use kube-logging namespace it cannot run ES pods and this error occurred but if I apply it on default namespace it works correctly.

how can I fix it and have the pods in logging name-space?

I am using single node cluster. I test it on other service provider's single node cluster and minikube and Kind the problem still persist. I think I'm missing something here.

[Reply](#) [Report](#)

^ [stalinbritto](#) July 12, 2019

5 Hi Team,

I have followed and when I execute elasticsearch_statefulset.yaml, am getting below output and it's taking a long time.

It couldn't step forward further.

```
root@hostname:~/kube# kubectl rollout status sts/es-cluster --namespace kube-logging
Waiting for 3 pods to be ready...
```

[Reply](#) [Report](#)

^ [josesantos](#) October 24, 2019

2

[SCROLL TO TOP](#)

Me too.

Have you solved?

[Reply](#) [Report](#)

^ [dakadivya](#) July 31, 2019

0 Can u make a video how to store application logs(destination)to EFK (source) server.

[Reply](#) [Report](#)

^ [jain007neeraj](#) October 2, 2019

1 Awesome Explanation, thanks Hanif, I just have one query how can I configure the fluentd to fetch logs from default namespace, Since right now it's just reading from kube-logging namespace.

[Reply](#) [Report](#)

^ [alex-davis0319](#) April 1, 2021

0 Any solution to this?

[Reply](#) [Report](#)

^ [Amipb](#) October 7, 2019

0 Great tutorial, but I'm stuck with the same kind of problem stated above : the rollout hangs at "Waiting for 3 pods to be ready".

[Reply](#) [Report](#)

^ [josesantos](#) October 24, 2019

0 Me too.

Have you solved?

[Reply](#) [Report](#)

^ [Amipb](#) October 28, 2019

0 No, not yet.

[Reply](#) [Report](#)

^ [HasanGoz](#) October 25, 2019

0 Same problem but the output is like this;

Waiting for 3 pods to be ready...

Waiting for 2 pods to be ready...

Waiting for 3 pods to be ready...

Waiting for 2 pods to be ready...

and stuck like this

[Reply](#) [Report](#)

SCROLL TO TOP

[phucphanhong](#) February 14, 2021



I had this issue, it is because you have no enough Node to run this infrastructure.
Increase the nodes or mark the cluster as scalable.

[Reply](#) [Report](#)



[sysadmin57e164f7fa058bd00e](#) October 26, 2019

Hi Hanif,

Thanks for this useful tutorial, i am successfully deployed EFK stack.

But in Kibana Dashboard i can't discover all namespaces that located on my k8s cluster ?!

[Reply](#) [Report](#)



[contraboy9999](#) November 18, 2019

Thanks for the detailed article. I am looking into implementing this but I wanted to know the minimum amount of storage the cluster should have which I need to get started on this. I will be scaling this out after the POC though. Thanks!!

[Reply](#) [Report](#)



[contraboy9999](#) December 6, 2019

Hi,

Thankk you for the detailed article:

I am having an issue when executing the following command.

I just get the below output and it goes in hung state right here. Below is the output:

```
kubectrl rollout status sts/es-cluster --namespace=kube-logging
```

Output:

Waiting for 3 pods to be ready...

I have 1 master and 2 nodes in my cluster. Any help would be appreciated.

[Reply](#) [Report](#)



[vinx](#) April 20, 2020

Everything works but I had to make a change to use my nfs server as a storage class for elastic search, I always had the error message: "pod has unbound immediate .."
but this was simply a problem of delay in the creation of the pvc, the new attempt should have been solved.

The real problem with was the init container fix-permissions, my volume nfs is with *allsquash* and *rootsquash*, there is no need to change the owner with *choown* in the init container (command: ["sh", "-c", "chown -R 1000: 1000 / usr / share / elasticsearch / data"]), the owner is always nobody, and furthermore this command is blocked, the init container fails and pod not start. I have removed the "fix-permission" init container and everything is working properly.

[SCROLL TO TOP](#)

In your own storage class installation you have to pay attention to this aspect.

I hope it will be useful to others

Vinx

[Reply](#) [Report](#)

^ [shivamsinghal](#) May 7, 2020

0 Thanks you for the detailed article:

I am having an issue when executing the following command.

I just get the below output and it goes in hung state right here. Below is the output:

```
kubectl rollout status sts/es-cluster --namespace=kube-logging
```

Output:

Waiting for 3 pods to be ready...

[Reply](#) [Report](#)

^ [sasidharallnew](#) June 21, 2020

0 i am also facing same issue any luck in solving ?

[Reply](#) [Report](#)

^ [kacperpb96](#) July 24, 2020

0 describe statefulset of your elasticsearch deployment and make sure you don't have problem with privileged pod permissions. In my case I has to create SecurityAccount/ClusterRole/ClusterRoleBinding and attach it in specs

[Reply](#) [Report](#)

Load More Comments




This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

SCROLL TO TOP

**GET OUR BIWEEKLY NEWSLETTER**

Sign up for Infrastructure as a
Newsletter.

**HOLLIE'S HUB FOR GOOD**

Working on improving health and
education, reducing inequality,
and spurring economic growth?
We'd like to help.

**BECOME A CONTRIBUTOR**

You get paid; we donate to tech
nonprofits.

[Featured on Community](#) [Kubernetes Course](#) [Learn Python 3](#) [Machine Learning in Python](#)
[Getting started with Go](#) [Intro to Kubernetes](#)

[DigitalOcean Products](#) [Virtual Machines](#) [Managed Databases](#) [Managed Kubernetes](#) [Block Storage](#)
[Object Storage](#) [SCROLL TO TOP](#) [Load Balancers](#)

Welcome to the developer cloud

DigitalOcean makes it simple to launch in the cloud and scale up as you grow – whether you're running one virtual machine or ten thousand.

[Learn More](#)



© 2021 DigitalOcean, LLC. All rights reserved.

Company

- [About](#)
- [Leadership](#)
- [Blog](#)
- [Careers](#)
- [Partners](#)
- [Referral Program](#)
- [Press](#)
- [Legal](#)
- [Security & Trust Center](#)

Products

- [Pricing](#)
- [Products Overview](#)
- [Droplets](#)
- [Kubernetes](#)
- [Managed Databases](#)
- [Spaces](#)
- [Marketplace](#)

Community

- [Tutorials](#)
- [Q&A](#)
- [Tools and Integrations](#)
- [Tags](#)
- [Product Ideas](#)
- [Presentation Grants](#)

Contact

- [Get Support](#)
- [Trouble Signing In?](#)
- [Sales](#)
- [Report Abuse](#)
- [System Status](#)

[SCROLL TO TOP](#)

- [Load Balancers](#)
- [Block Storage](#)
- [API Documentation](#)
- [Documentation](#)
- [Release Notes](#)
- [Hatch Startup Program](#)
- [Shop Swag](#)
- [Research Program](#)
- [Open Source](#)
- [Code of Conduct](#)

SCROLL TO TOP